



TensorFlow Multilayer Perceptron Modeling for Heart Disease Prediction

William Lorenzo

California State University - East Bay

STAT-312-02-1

Dr. Taejoon Kim

03/16/2026

Abstract

Cardiovascular disease remains as one of the primary sources of impermanence worldwide, portraying early detection and predictive modeling as a crucial attribute in modern healthcare analytics. This examination applies a feedforward neural network, particularly a Multilayer Perceptron (MLP), integrated in TensorFlow to predict the presence of heart disease using clinical and demographic features. Statistical practice such as exploratory data analysis (EDA) was carefully conducted to grasp a solid foundation of the distribution of the dataset and linear relationships among certain features. Qualitative variables were cautiously encoded using one-hot encoding. Quantitative predictors were standardized to refine neural network conjunction. Two feedforward neural network models were constructed, a baseline and a tuned neural network to utilize regularization techniques such as L2 weight penalties, batch normalization, dropout, learning rate scheduling, and early stopping. Model performance was assessed using the traditional metrics for classification tasks, namely accuracy, precision, recall, F1-score, and ROC-AUC. Results illustrate that the tuned model attained improved predictive performance and generalization when under the comparison of the baseline feedforward neural network model. These matters of course attentively display the true efficacy of neural network models for tabular-like medical classification tasks, notably binary classification.

Introduction

Heart disease remains as one of the most notorious public health apprehensions worldwide, taking account for more than a plethora of deaths each year (World Health Organization, 2023). Early detection and predictive modeling of cardiovascular conditions are necessary for enhancing medical results and enabling preventative care strategies. Machine learning algorithms have increasingly been applied to clinical datasets in order to distinguish complex relationships between patient characteristics and disease outcomes (Rajkomar et al., 2019).

Artificial neural networks, especially the Multilayer Perceptrons (MLPs), are capable of modeling nonlinear relationships between features and target variables. These models comprise various layers of interconnected neurons that learn hierarchical representations of input features through iterative optimizations (Goodfellow et al., 2016). In the context of medical classification procedures, neural networks have the ability to grasp precise interactions among physiological features that conventional machine learning models may be portrayed as ineffective.

This practice consists of a TensorFlow-based feedforward neural network to predict whether an individual has heart disease using a dataset comprising of demographic and clinical features. The study proceeds through numerous phases including exploratory data analysis (EDA), feature engineering, model development, and performance evaluation. Two feedforward neural network models are crafted, a baseline neural network using all available features and a tuned variant containing regularization and training optimization procedures. The intention of this study is to carefully evaluate whether tuning improves predictive accuracy and generalization upon identifying individuals with heart disease.

Theoretical Background

Artificial Neural Networks

Artificial neural networks are computational models which were previously inspired by the neurological structures of the brain. A feedforward neural network consists of an input layer, one or more hidden layers, and an output layer. Each neuron computes a weighted sum of its inputs, which is then proceeded by an activation parameter. In statistical terms, the output of a single neuron can be conveyed as

$$z = w^T x + b$$

where

$$w = (w_1, w_2, \dots, w_n)$$

represents the weight vector, x represents the input vector, and b is the bias term.

The neuron output is then transformed by an activation function:

$$a = \sigma(z)$$

where $\sigma(\cdot)$ illustrates the activation function.

In this process, the Rectified Linear Unit (ReLU) activation function is used in hidden layers:

$$ReLU(z) = \max(0, z)$$

The output layer utilizes a sigmoid activation function to create probabilities for binary classification:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

The sigmoid function forces predictions to the interval $[0, 1]$, allowing them to be interpreted as probabilities of heart disease.

Loss Function

For binary classification procedures, the binary cross-entropy loss function is frequently utilized:

$$L(y, \hat{y}) = -\frac{1}{n} \sum_{i=1}^n [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

Where y_i shows the true label and $\hat{y}_i$ represents the predicted probability. The objective is to minimize this loss function as much as possible.

Gradient Descent Optimization

Training neural networks involves minimizing the loss functions with respect to model parameters using gradient-based optimization. Gradient descent updates weights according to:

$$w_{t+1} = w_t - \eta \nabla L(w_t)$$

Where η is the learning rate and $\nabla L(w_t)$ represents the gradient of the loss function. This study uses the Adam optimizer, which combines momentum and adaptive learning rates to further advance convergence (Kingma & Ba, 2015).

Regularization Techniques

To improve generalization and prevent overfitting, various regularization procedures were carefully integrated.

L2 Regularization

L2 regularization penalizes large weights upon adding a penalty term:

$$L_{reg} = L + \lambda \sum w^2$$

This discourages excessively large parameter values.

Dropout

Dropout randomly disables a proportion of neurons in duration of the training procedure:

$$h_i^{(\text{drop})} = \begin{cases} 0 & \text{with probability } p \\ h_i & \text{with probability } 1 - p \end{cases}$$

This prevents neurons from co-adapting and improves generalization (Srivastava et al., 2014).

Batch Normalization

Batch normalization standardizes intermediate activations:

$$\hat{x} = \frac{x - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

This stabilizes training and allows for higher learning rates (Ioffe & Szegedy, 2015)

Exploratory Data Analysis

The dataset consists of 918 observations and multiple clinical variables related to cardiovascular health. Exploratory data analysis was performed to grasp how variables were distributed and the potential relationships with the target variable while assessing class balance. Out of the 918 individuals in the dataset, 508 individuals have heart disease while the remaining 410 do not have heart disease.

Correlation analysis was conducted to further examine linear correlations between features and the target variable. Weak to moderate correlations were studied among several predictors, such as age, fasting blood sugar, and exercises-related variables. Features such as maximum heart rate and ST depression portrayed a strong association with heart disease outcomes.

Categorical variables were converted into dummy variable using the one-hot encoding procedure in order to incorporate them in correlation analysis and model training.

Methodology

Data Preprocessing

Before training the neural network architecture, various preprocessing steps were thoroughly executed, one-hot encoding for qualitative features, feature scaling using standardization, and train-test-split where the dataset is divided into 70% training and 30% testing. Standardization transforms each predictor as

$$x' = \frac{x - \mu}{\sigma}$$

Where μ is the mean and σ is the standard deviation.

Feature scaling is crucial when it comes to neural networks, this ensures that all input variables contribute proportionally to gradient revisions.

Baseline Neural Network Architecture

The baseline model consisted of a simple feedforward neural network with two hidden layers. Architecture possessed 15 features in the input layer, 64 and 32 neurons in the dense layers (with a ReLU activation), and 1 neuron for the output layer (utilizing the sigmoid activation). The model was trained using the Adam optimizer and binary cross-entropy loss. These results indicate that the baseline model achieved a training accuracy of approximately 0.9330, test accuracy of 0.8696, with a test loss of 0.3499. These results indicate that the baseline model correctly classifies approximately 93.3% of the observations in the dataset. The baseline model learned patterns specific to the training but did not generalize effectively to unseen data, which suggests overfitting due to a higher training accuracy when compared to testing accuracy.

Tuned Neural Network Architecture

In order to better-enhance generalization, a tuned variant of the neural network model was crafted. After applying numerous model tuning techniques, including L2 regularization,

batch normalization, dropout layers, learning rate scheduling, and early stopping, the tuned neural network achieved a training accuracy of 0.8816 and a test accuracy of 0.8768, with a test loss of 0.3410. Dense layers kept the 64-neurons configuration with ReLU activation and L2 regularization, implementation of batch normalization, dropout set to 0.3, 32-neurons with ReLU activation and L2 regularization, batch normalization, dropout adjusted to 0.2, and an output layer with a sigmoid activation. The training process was monitored using validation loss and early stopping was applied to prevent excessive training once performance plateaued.

Model Evaluation

Model performance was evaluated using traditional classification metrics:

Accuracy:

$$Accuracy = \frac{True\ Positives + True\ Negatives}{True\ Positives + True\ Negatives + False\ Positives + False\ Negatives}$$

Precision:

$$Precision = \frac{True\ Positives}{True\ Positives + False\ Positives}$$

Recall:

$$Recall = \frac{True\ Positives}{True\ Positives + False\ Negatives}$$

F1 Score:

$$F1 = 2 \times \frac{Precision \times Recall}{Precision + Recall}$$

ROC-AUC measures the area under the receiver operating characteristic curve, representing the model's ability to differentiate between certain classes.

The tuned feedforward neural network model achieved improved performance across all evaluation metrics when under the comparison to the baseline model.

Results

The tuned neural network achieved a higher predictive performance across multiple evaluation metrics such as accuracy, precision, recall, F1-score, and ROC-AUC. This indicates that the combination of regularization techniques, architectural refinement, and feature selection heavily corrected the model's ability to generalize to unseen data. The ROC-AUC score above 0.95 suggests that the tuned model possesses strong discriminative ability for identifying individuals with heart disease.

Discussion

The results demonstrate that neural network tuning techniques can dramatically improve predictive performance in medical classification tasks. Regularization procedures such as dropout and L2 penalties reduce overfitting by limiting the complexity of the architecture, while batch normalizing stabilizes training dynamics. Additionally, adaptive optimizers such as Adam allow efficient training of deep neural networks by adjusting learning rates during optimization.

Despite the absolute performance of the tuned model, neural networks applied to tabular datasets often face various limitations when under the comparison to specialized machine learning algorithms, not to mention gradient boosting, transformers, support vector machine (SVM), etc. In terms of future studies, ensemble methods and feature engineering strategies would be considered to further enhance predictive performance.

Model Evaluation

This practice applied TensorFlow-based neural networks to predict heart disease while integrating clinical data. A baseline Multilayer Perceptron model was primarily developed and eventually refined with the assistance of regularization, hyperparameter tuning, and various model adjustments. The tuned model demonstrated phenomenal predictive performance across

various classification metrics, which indicates an improved generalization and reduced the presence of overfitting. These results highlight the effectiveness of neural network models for healthcare analytics.

References

- Ian Goodfellow, Y. B. (2016). *Deep Learning* (1st ed.). The MIT Press.
- Ioffe, S., & Szegedy, C. (2015). Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift. *arXiv.Org*.
- Kingma, D. P., & Ba, J. (2017). Adam: A Method for Stochastic Optimization. *arXiv.Org*.
- Machine Learning in Medicine. (n.d.). *The New England Journal of Medicine*.
<https://doi.org/10.1056/nejmc1906060>
- Hinton, G. E., Srivastava, N., Krizhevsky, A., Sutskever, I., & Salakhutdinov, R. R. (2012). Improving neural networks by preventing co-adaptation of feature detectors. *arXiv.Org*.
- World Health Organization. (n.d.). *Cardiovascular diseases (cvds)*. World Health Organization.
[https://www.who.int/news-room/fact-sheets/detail/cardiovascular-diseases-\(cvds\)](https://www.who.int/news-room/fact-sheets/detail/cardiovascular-diseases-(cvds))